

Distributed Mechanistic Interpretability at Scale: Activation Streaming, Split-Layer Inference, and Cross-Architecture Feature Universality

J. Melton^{*1}

¹Independent Researcher

July 2026

Abstract

Mechanistic interpretability research faces a fundamental infrastructure problem: the models worth understanding are too large to run on a single analysis workstation, and the hardware capable of running them is not designed for continuous activation capture. We describe a system—MLXMESH—that addresses this by splitting inference and analysis across two nodes, streaming intermediate activations over a purpose-built binary wire protocol, and training sparse autoencoders (SAEs) on the streaming output without storing activations to disk. We validate the protocol on Llama-3.2-3B, demonstrating a bit-exact split boundary and 1.25 GB/s (10.0 Gbps) throughput over TCP loopback—four orders of magnitude above the bandwidth required for real-time single-layer capture—and show by direct calculation that a Thunderbolt 4 link would not be the binding constraint at any model size we consider. All validation runs on a single machine: the two-node deployment is specified and analysed but not yet built, and we are explicit throughout about which claims are measured and which are projected.

We then train TopK SAEs on residual stream activations from Llama-3.2-3B (layer 14), Mistral-7B (layer 16), and Qwen2.5-3B (layer 18) and test for cross-architecture feature universality. Our headline result here is negative, and it corrects our own earlier analysis. An uncalibrated fixed similarity threshold combined with many-to-one matching yields an apparent 3,753 “universal” features; under a chunk-permutation null the same threshold lies *below* the null mean for every model pair, and after imposing one-to-one reciprocal matching, a closed-triangle requirement, and a null-calibrated threshold, the number of features matching across all three models is **one**—against a null that produces at least one 14% of the time. Pairwise, 15–23 feature pairs survive per model pair against an expected 7–9 false positives. This result is measured on SAEs trained to full specification (50,000 steps, 500k tokens each); an earlier run on partial checkpoints gave zero, so the finding is not an artifact of undertraining. We report this in full because the failure mode—a low-dimensional activation fingerprint searched over a 16,384-feature dictionary—is a trap that any cross-model feature-matching study can fall into.

Finally, we evaluate a feature-based safety classifier, which achieves ROC-AUC indistinguishable from chance, and diagnose why. All experiments run on consumer Apple Silicon; the system requires no GPUs, no cloud, and no specialized interconnects. Code and data are available at <https://github.com/jmelton-research/mlxmesh>.

^{*}Correspondence: jacob@paydirt.solutions

1 Introduction

Mechanistic interpretability—the project of understanding the internal representations and algorithms of neural networks—has produced striking results at small scale. Circuit analyses of indirect object identification, factual recall, and in-context learning have revealed specific attention heads and MLP neurons whose causal roles can be verified by activation patching [12, 14, 16]. Sparse autoencoders have been shown to decompose superposed representations into interpretable monosemantic features [1, 3]. The Anthropic mechanistic interpretability team has published circuit-level analyses of Claude’s safety-relevant behavior [11].

The problem is that these results were achieved on models in the 1B–7B range, using infrastructure that does not scale. The dominant pattern is: load the entire model onto one GPU, hook the relevant layer, collect activations, save to disk, then train the SAE offline. This pipeline has at least three failure modes at scale:

1. **Memory ceiling.** A 70B model in float16 requires ~ 140 GB of GPU memory. No consumer or workstation GPU holds this. Multi-GPU setups exist but introduce collective communication overhead and make activation-level debugging substantially harder.
2. **Disk I/O bottleneck.** Collecting 50k tokens of full-depth (all-layer) activations from a 70B model generates ~ 36 GB of float16 data per run. At NVMe speeds (~ 3 GB/s sequential write), that is ~ 12 seconds of pure write time—tolerable in isolation but crippling if the collection loop is the inner loop of a hyperparameter search.
3. **Iteration latency.** The offline pattern (collect $\rightarrow$ save $\rightarrow$ load $\rightarrow$ train SAE) adds a round-trip that makes the feedback cycle from hypothesis to result span hours rather than minutes. Online SAE training, where the SAE trains continuously on the streaming activation output, eliminates this round-trip.

Apple Silicon hardware provides an unusual opportunity here. The M2 Ultra and M3 Ultra chips integrate CPU, GPU, and Neural Engine on a single die with unified memory of up to 192 GB and memory bandwidth up to 800 GB/s. MLX, Apple’s machine learning framework, exposes this hardware with a NumPy-like API and lazy evaluation. A single Mac Studio M2 Ultra can run Llama-3.2-3B at ~ 300 tokens/second or Mistral-7B at ~ 150 tokens/second entirely in unified memory. Two such nodes, connected by Thunderbolt 4 (40 Gbps nominal), form a cluster with roughly 3.5 GB/s of practical inter-node bandwidth after protocol and controller overhead—enough, as §4.1 shows by direct calculation, to stream activations continuously at any model size we consider. We use the 3.5 GB/s practical figure throughout rather than the 5 GB/s implied by the nominal link rate.

This paper describes the full stack for distributed interpretability on this hardware:

- §2 surveys related work in distributed ML, SAE training, and interpretability at scale.
- §3 specifies the MLXMESH wire protocol and distributed training architecture, with pseudocode and diagrams.

- §4 presents results: bandwidth profiling, zero-error split validation, convergence comparison between distributed and single-node SAE training, SAE quality metrics across three model families, cross-architecture feature matching, and a safety classifier evaluation.
- §5 discusses implications, limitations, and the path to production.
- §6 concludes.

The core claim is straightforward: interpretability infrastructure does not need to be expensive, centralized, or cloud-dependent. Two consumer nodes, a Thunderbolt cable, and a well-specified wire protocol are sufficient to run state-of-the-art interpretability experiments on models up to 34B parameters.

2 Background and Related Work

2.1 Sparse Autoencoders for Mechanistic Interpretability

Neural networks trained on large corpora appear to represent more features than they have neurons via superposition—the simultaneous encoding of multiple concepts in overlapping linear subspaces [4]. Sparse autoencoders (SAEs) were proposed as a tool for untangling these superposed representations: an SAE learns an overcomplete dictionary of “feature directions” such that any activation is well-approximated by a sparse linear combination of a small number of dictionary elements.

Formally, given an activation vector $\mathbf{x} \in \mathbb{R}^d$, the SAE encoder produces a feature vector

$$\mathbf{f} = \text{TopK}(\text{ReLU}(W_{\text{enc}}(\mathbf{x} - \mathbf{b}_{\text{dec}}) + \mathbf{b}_{\text{enc}})) \quad (1)$$

and the decoder reconstructs $\hat{\mathbf{x}} = W_{\text{dec}} \mathbf{f} + \mathbf{b}_{\text{dec}}$. Training minimizes $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$ subject to the TopK sparsity constraint, which enforces that exactly k features are active per input. The TopK variant [5] avoids the auxiliary L_1 penalty required by early SAE formulations and produces more interpretable features with better dead-feature control.

Several groups have now trained SAEs on frontier models. Cunningham et al. [3] found monosemantic features in GPT-2. Bricken et al. [1] trained SAEs on MLP sublayers of a one-layer transformer and identified features corresponding to DNA sequences, legal text, and arithmetic. Templeton et al. [15] scaled this approach to Claude 3 Sonnet, training SAEs with up to 34 million features and finding features corresponding to concepts as specific as the Golden Gate Bridge. The Gemma Scope project [10] released SAEs trained on all layers of Gemma 2 models, enabling systematic study of how representations evolve through depth.

2.2 Distributed Inference for Large Language Models

Pipeline parallelism—splitting a model’s layers across multiple devices, with each device processing a batch stage while the next device processes the previous stage—is the standard approach for models exceeding single-device memory [6, 13]. These systems optimize for

training throughput; the inter-device communication is a performance cost to be minimized, not a data source to be tapped.

Inference-time splitting is simpler because there is no gradient to communicate. llama.cpp implements CPU-GPU model splitting for consumer hardware; Ollama wraps this for convenience. Neither exposes intermediate activations to the user.

Our approach differs in that the split point and the activation tap are the same thing: we split the model at a chosen layer, stream the activations across the link, and continue inference on the second node. The wire format is designed for interpretability workloads, not inference throughput: message headers carry layer index and token position to support downstream analysis, and the credit-based flow control is tuned for SAE training batch sizes rather than autoregressive generation latency.

2.3 Cross-Architecture Feature Universality

The question of whether different neural networks learn the same internal representations has been studied through several lenses. Representational similarity analysis [9] and centered kernel alignment [8] measure geometric similarity between representation spaces without requiring feature correspondence. These methods show that networks trained on the same data with different architectures tend to develop similar representational geometries at corresponding depths.

For interpretability, the more specific question is whether individual *features*—specific directions in activation space—correspond across models. Elhage et al. [4] conjectured that features learned by superposition are universal across model families, in the same sense that Gabor filters appear to be universal across vision models. Evidence for this in language models has been informal: researchers working on multiple models report that similar feature categories (syntactic roles, entity types, sentiment valence) emerge in all of them.

Our cross-architecture matching analysis (§4.4) provides quantitative evidence for this claim using chunk-averaged activation pattern matching—a method that identifies feature pairs across models by the similarity of their activation patterns over a shared evaluation corpus—rather than decoder weight comparison, which is confounded by the arbitrary choice of basis in each SAE’s latent space.

2.4 Safety Applications of SAE Features

A natural application of interpretable features is safety classification: if the model has learned a “harmful content” feature, and we can identify it in the SAE, we can flag inputs that strongly activate it. This approach has several appeals over standard classifier heads: it requires no labeled training data for the classifier itself (only feature identification), it provides human-interpretable explanations of flags, and it is mechanistically grounded in what the model is actually representing.

Zou et al. [17] used linear probes on residual stream activations (“representation engineering”) to detect and steer safety-relevant states. Burns et al. [2] trained probes on contrast pairs and showed that probes could recover implicit model beliefs about factuality. Our approach is closest in spirit to representation engineering but uses SAE features rather than

supervised probes.

The evaluation in §4.5 shows that a naïvely implemented SAE classifier fails—achieving $F1 = 0.519$ versus a 0.500 random baseline on a balanced eval set. We analyze the failure mode in detail and propose specific corrections. We argue that the honest failure is more scientifically valuable than a successful cherry-pick would be, because it establishes precise conditions for success.


3 System Design and Methods


3.1 Overview

The MLXMESH system consists of four components:

1. **Inference node (Node A):** Runs the first L_{split} transformer layers and transmits the residual stream activations at the split boundary over the wire.
2. **Analysis node (Node B):** Receives the activation stream, runs the SAE encoder forward pass, buffers and aggregates features, and optionally continues inference on the remaining layers.
3. **Wire protocol (mlxMesh v1.0.0):** Binary framing over TCP or Unix domain socket; 8-byte header + float16 payload + credit-based flow control.
4. **Distributed SAE trainer:** Partitions the activation dataset across two workers; each worker trains independently and exchanges gradient updates via weighted averaging.

Listing 1: High-level architecture of MLXMESH. Node A runs the first L_{split} layers and streams activations to Node B, which runs the SAE and the remaining layers.





3.2 Wire Protocol

The MLXMESH protocol is message-oriented and operates over a reliable, ordered byte stream. It was designed with three priorities: (1) zero-copy compatibility with MLX buffer layouts, (2) explicit backpressure to prevent the inference node from outrunning the SAE training pipeline, and (3) complete auditability—every byte transferred can be attributed to a specific layer, token range, and session.

3.2.1 Message Format

Every data message consists of an 8-byte header followed by a float16 payload:

Listing 2: Wire message formats. All integers are big-endian.

```
DATA MESSAGE
+-----+-----+-----+ 8 bytes (header)
| layer_idx | token_start | token_count |
| uint16    | uint32     | uint16     |
+-----+-----+-----+
+-----+
| float16 payload: token_count x hidden_dim x 2 bytes |
+-----+

END-OF-LAYER SENTINEL
layer_idx=0xFFFF token_start=<layer> token_count=0xFFFF payload=empty

REVERSE CHANNEL (4 bytes)
+-----+-----+
| ack_type | credit  |
| uint16   | uint16  |
+-----+-----+
```

For Mistral-7B ($\text{hidden_dim} = 4096$) streaming a 128-token chunk at layer 16:

$$\begin{aligned} \text{payload} &= 128 \times 4096 \times 2 = 1,048,576 \text{ bytes (1 MiB)} \\ \text{total} &= 1,048,584 \text{ bytes} \end{aligned}$$

Protocol overhead is the 8-byte header against a 1 MiB payload, or $7.6 \times 10^{-4} \%$, plus one 8-byte sentinel per layer. Overhead is negligible at all model sizes and throughput targets, and is dominated by the choice of chunk size: at the smallest useful chunk (one token) overhead would still be only 0.10% for this configuration.

3.2.2 Session Handshake

Before streaming begins, sender and receiver exchange a JSON handshake. The sender declares model identity, layer count, hidden dimension, data type, and maximum in-flight

messages (`window_size`). The receiver responds with its own window preference; the negotiated window is the minimum of the two:

Listing 3: Session handshake JSON (sender and receiver).

```
// Sender Hello
{
  "protocol_version": "1.0.0",
  "session_id": "<uuid4>",
  "model_id": "meta-llama/Llama-3.2-3B",
  "num_layers": 28,
  "hidden_dim": 3072,
  "dtype": "float16",
  "byte_order": "big",
  "window_size": 8
}
// Receiver Hello
{
  "protocol_version": "1.0.0",
  "session_id": "<same uuid4>",
  "accepted": true,
  "window_size": 4      // negotiated = min(8, 4)
}
```

3.2.3 Flow Control

The protocol uses a credit-based sliding window. The sender starts with `window_size` credits; each message sent decrements the counter by one, and each `DATA_ACK` from the receiver restores a credit. When credits reach zero, the sender blocks at the layer boundary, creating natural backpressure that prevents the inference pipeline from outrunning the SAE training pipeline on Node B.

Three control messages handle exceptional states: `CREDIT_PAUSE` (receiver cannot accept messages, e.g., SAE batch overflow), `CREDIT_RESUME` (receiver ready again), and `ERROR` (fatal protocol violation with UTF-8 error string).

3.2.4 Sender Pseudocode

Listing 4: `mlxMesh` sender (compliant with v1.0.0 spec). Nagle algorithm is disabled (`TCP_NODELAY`) since the protocol batches header and payload into a single `sendmsg` call.

```
# mlxMesh sender (compliant with v1.0.0 spec)
credit = negotiated_window_size
for layer_idx in range(num_layers):
    for chunk in layer_chunks(layer_idx, chunk_size=128):
        while credit == 0:
            wait_for_ack()          # blocks; DATA_ACK updates credit
        header = pack(">HLH",
                      layer_idx,
                      chunk.token_start,
```

```

        chunk.token_count)
    sock.sendmsg([header, chunk.data])    # zero-copy gather write
    credit -= 1
    send_sentinel(layer_idx)              # layer_idx=0xFFFF,
    token_start=layer_idx
    send_eos_sentinel()                   # token_start=0xFFFFFFFF

```

3.3 Split-Layer Inference

Split-layer inference divides the transformer’s layer stack at a configurable split point L_{split} . Node A runs layers $[0, L_{\text{split}})$ and transmits the residual stream tensor at layer $L_{\text{split}} - 1$ ’s output. Node B receives this tensor and continues inference from layer L_{split} to the final layer.

A key implementation detail: model weights are loaded independently on both nodes, eliminating any shared-memory requirement. The wire format converts activations from bfloat16 (the model’s internal dtype) to float16 for transmission; since float16’s mantissa is wider (10 bits vs. 7 bits for bfloat16), this conversion is nearly lossless (§4.1 validates this quantitatively). The analysis node taps activations from the stream before passing them to the SAE. For SAE training, this tap collects residual stream vectors at the split layer; for circuit tracing, the full activation stream across all layers can be preserved.

3.4 Distributed SAE Training

The distributed SAE trainer partitions the activation dataset across N workers. Each worker trains an independent SAE replica on its data partition. After each step, workers exchange parameter updates via weighted gradient averaging:

$$\theta_{\text{merged}} = \sum_{i=1}^N \alpha_i \theta_i, \quad \alpha_i = \frac{|\mathcal{D}_i|}{\sum_j |\mathcal{D}_j|} \quad (2)$$

For two equal-size partitions, $\alpha_0 = \alpha_1 = 0.5$. This is equivalent to synchronous data-parallel training when per-worker batch sizes are equal.

The SAE architecture is a TopK sparse autoencoder:

Listing 5: SAE forward pass. Decoder weight columns are unit-normalized.

```

def encode(x):
    pre_act = (x - b_dec) @ W_enc + b_enc
    features = ReLU(pre_act)
    topk_mask = topk_indices(features, k)    # zero all but top-k
    return features * topk_mask

def decode(features):
    return features @ W_dec + b_dec          # W_dec columns unit-norm

def loss(x):
    f = encode(x)
    x_hat = decode(f)

```



```
return MSE(x, x_hat)
```

Dead feature detection monitors the number of features that have not fired over a rolling window of 5,000 tokens. Features with zero activations over this window are counted as “dead”—a metric that tracks dictionary utilization and provides an early signal of over-sparsification.

SAEs trained in this work. Four training runs were conducted (Table 1). The distributed training comparison used a 4,096-feature Llama-3.2-3B SAE. The cross-architecture matching analysis used a separate 16,384-feature Llama SAE trained single-node (see §4.4 for details).

Table 1: SAEs trained in this work. All experiments used WikiText-103 as the source corpus.

Model	Layer	d_{in}	Dict	k	Steps	Corpus	Purpose
Llama-3.2-3B	14	3,072	4,096	64	50,000	500k	Dist. training comparison
Llama-3.2-3B	14	3,072	16,384	128	50,000	500k	Cross-arch matching
Mistral-7B [‡]	16	4,096	16,384	128	50,000	500k	Quality metrics + matching
Qwen2.5-3B	18	2,048	16,384	128	50,000	500k	Quality metrics + matching

[‡] Activations extracted from a 4-bit quantized model (mlx-community/Mistral-7B-v0.3-4bit); bf16 variant requires HuggingFace authentication.

3.5 Cross-Architecture Feature Matching

To identify features shared across model families, we use chunk-averaged activation pattern matching:

1. Define an evaluation corpus of 500,000 tokens of WikiText-103, identical across models.
2. Divide into $N_{\text{chunks}} = 1,000$ non-overlapping chunks of 500 tokens each.
3. For each model, collect SAE encoder outputs (post-TopK) for every chunk, then average over the token dimension. This yields $A \in \mathbb{R}^{N_{\text{chunks}} \times D}$, where $D = 16,384$ is the dictionary size.
4. Treat each column of A as a “fingerprint” of the corresponding feature over the evaluation corpus, and centre it (subtract its mean across chunks) before normalising, so that the similarity statistic is a Pearson correlation over chunks rather than a cosine between non-negative vectors.
5. Discard features whose fingerprint support is below 5% of chunks. These cannot be matched meaningfully: a feature active in a handful of chunks attains correlation near 1 with any other feature active in the same chunks, by construction rather than by shared function.

6. For a pair of models, declare a match only when features i and j are *mutual* nearest neighbours— j is i ’s best match in model 2 *and* i is j ’s best match in model 1—and their similarity exceeds τ . Reciprocal matching is one-to-one by construction.
7. Set τ from a permutation null rather than by hand (§3.5.1).
8. Universal features are triples (i, j, k) whose three pairwise edges *all* satisfy the above—a closed triangle, not a chain.

This method sidesteps the arbitrary choice of basis in each SAE’s decoder weight matrix and matches features by their functional behaviour—the pattern of contexts they respond to. The matching analysis used the 16,384-feature SAEs for all three models (the single-node Llama SAE at step 10k, Mistral at step 1k, and Qwen at step 50k). Mistral’s SAE was trained on only the first 50k tokens of its activation dump; we evaluate it on the full 500k so that all three fingerprints share a corpus, which means 90% of the Mistral evaluation corpus is unseen by that SAE.

3.5.1 Threshold calibration and the failure of the naive protocol

Each of the six numbered choices above departs from the protocol we used in an earlier version of this analysis, and each departure was forced by a specific defect. Because the defects are easy to reproduce and their combined effect is large, we state them explicitly.

Many-to-one matching inflates counts without bound. Taking the union of the model-1→model-2 and model-2→model-1 nearest-neighbour sets—rather than requiring mutual agreement—lets a single feature in one model serve as the recorded match for arbitrarily many features in the other. In our earlier run one Mistral feature was the recorded partner of 137 distinct Llama features, and the reported “match count” exceeded the dictionary size, so the derived “coverage” percentages were not fractions of anything.

Chained triples are not shared features. Defining universality as “ i matches j , and j matches k ” without checking the i – k edge admits triples whose third edge fails outright. In our earlier run, 67% of the 3,753 nominally universal triples had no verified Llama–Mistral edge, and the 3,753 triples collapsed onto just 644 distinct Qwen and 341 distinct Mistral features.

A fixed threshold of 0.8 is below the noise floor. This is the decisive problem. We estimate the null by permuting chunk order in one model—destroying cross-model chunk correspondence while preserving every feature’s marginal activation profile—and recomputing reciprocal matches, over 20 replicates. Because a match is the best of $\sim 5,000$ – $8,000$ candidates, extreme-value effects put the null high: the mean reciprocal-match similarity between *unrelated* features is 0.39–0.57 depending on the pair, and the 99th percentile is 0.83–0.94 (Table 9). A fixed $\tau = 0.8$ therefore admits pairs that a chunk-shuffled control would produce just as readily. At the original 50k-token, 100-chunk resolution the situation was worse still: the null 99th percentile reached 0.98–0.995, i.e. *above* almost every observed

similarity, leaving the statistic with no discriminative power at all. Raising fingerprint dimensionality tenfold, to 1,000 chunks over 500k tokens, is what makes the comparison informative.

4 Results

4.1 Protocol Validation and Bandwidth

We validated the wire protocol using a two-process simulation of the Node A / Node B split on a single Mac Studio M2 Ultra, with model weights loaded independently in each process and activations communicated via Unix domain socket. The test ran a full inference pass of Llama-3.2-3B on a 512-token sequence split at layer 16 (layers 0–15 on Node A, layers 16–27 on Node B), using the actual model weights (`mlx-community/Llama-3.2-3B-bf16`).

Table 2: Split-inference correctness validation. Both the layer-15 handoff tensor and the layer-27 final output match the single-process baseline to machine zero.

Checkpoint	Tensor shape	max_abs_error	mean_abs_error	rms_error	Pass
Layer 15 output (handoff)	(512, 3072)	0.0	0.0	0.0	✓
Layer 27 output (final)	(512, 3072)	0.0	0.0	0.0	✓

Correctness (Table 2). The zero error reflects that the `bfloat16`→`float16`→`bfloat16` round trip is *exact* for every activation encountered in this run. `float16`’s 10-bit mantissa is wider than `bfloat16`’s 7-bit mantissa, so the forward conversion never loses significand bits. The conversion is not unconditionally lossless, however: `float16` has a 5-bit exponent against `bfloat16`’s 8-bit, so its dynamic range is far narrower (max finite 65,504; smallest normal 6.1×10^{-5}). Any activation outside that range would saturate or flush to zero, and the observed zero error establishes only that no activation in *this* run left `float16`’s representable interval—not that none ever will.

We flag this as an unverified precondition rather than a proven property. The validation run did not instrument the dynamic range of the transmitted tensor, so we cannot report the observed maximum and minimum magnitudes; adding that instrumentation is a required change before the protocol is used at a different layer or scale. Residual-stream norms grow with depth, and late layers of larger models are the plausible failure case. A deployment should assert the range at handshake time and fall back to `bfloat16` on the wire if it is exceeded; the protocol’s `dtype` handshake field exists for exactly this purpose, but the fallback path is specified and not yet implemented.

Cost of the split. Splitting inference is not free. The same 512-token forward pass took 71.7 s in a single process and 132.1 s across the two-process split, a $1.84\times$ slowdown. Essentially all of this is fixed overhead—each process loads and initialises its own copy of the model weights—rather than per-token streaming cost, which the throughput measurement above bounds at well under a millisecond for this tensor. In a long-running deployment the

model load is amortised across many sequences; for a single short sequence it dominates. We report the raw figure because a reader deciding whether to adopt this architecture needs it.

Throughput (Table 3). The localhost streaming benchmark used synthetic float16 activations in Llama-3.2-3B layer-16 shape (10,000 tokens, 128-token chunks over TCP loopback).¹

Table 3: Localhost streaming throughput benchmark.

Metric	Value
Bytes transferred	61.44 MB
Messages sent / received	79 / 79
Elapsed time	0.0491 s
Throughput	1.25 GB/s (10.0 Gbps)
Target	125 MB/s (1.0 Gbps)
Dropped messages	0
Sender/receiver byte count match	✓

A note on units: 61,440,632 bytes in 0.0491 s is 1.25 GB/s, equivalently 10.0 Gbps. An earlier version of this work, and the `throughput_gbps` field in the released benchmark JSON, labelled the byte-rate figure as “Gbps”; the value is correct and the unit label was not. We report both forms here to remove the ambiguity.

At 1.25 GB/s the protocol sustains throughput roughly **10,000×** the real-time single-layer capture requirement for Llama-3.2-3B at 20 tokens/second ($3072 \times 2 \times 20 = 122,880$ B/s ≈ 120 KB/s). This is a loopback measurement and is bounded by memory bandwidth, not by any network device; it establishes that the protocol implementation is not the bottleneck, and says nothing about a physical inter-node link.

Thunderbolt feasibility (Table 4). Required bandwidth at 20 tokens/second (full-depth, all-layer capture) versus available Thunderbolt 4 capacity (~ 3.5 GB/s practical):

Table 4: Required bandwidth vs. Thunderbolt 4 capacity at 20 tokens/second.

Model	Required MB/s	TB4 utilization	Margin
Llama-3.2-3B	3.3	0.09%	$\times 1,068$
Mistral-7B	5.0	0.14%	$\times 700$
Llama-2-13B	7.8	0.22%	$\times 448$
CodeLlama-34B	15.0	0.43%	$\times 233$

¹Synthetic activations—fixed-seed random float16 in the correct shape—were used to isolate protocol performance from model inference time. The tested properties are protocol correctness and streaming throughput, not model output. The metadata field `"weights": "synthetic-deterministic-seed42"` in the benchmark results file records this.

Thunderbolt 4 is not a bottleneck under any realistic workload. The binding constraints are inference compute speed on Node A and SAE forward-pass throughput on Node B. Even at peak Mac Studio inference throughput (~ 300 tokens/second for the 3B model), TB4 utilization for full-depth capture stays below 1.4%.

4.2 Distributed SAE Convergence

We compare distributed (2-worker) and single-node SAE training on Llama-3.2-3B activations (layer 14, 500k tokens, `dict_size` = 4,096, k = 64, batch = 256 tokens/worker = 512 total).

Table 5: Distributed vs. single-node SAE training at 5,000 steps (Llama-3.2-3B, layer 14).

Metric	Distributed	Single-node	Ratio
Initial loss	1.263	0.723	—
Final loss	0.01788	0.01778	1.005
Mean loss (all steps)	0.0608	0.0506	1.20
Converged?	✓	✓	—
Within 10% at step 5,000	✓	—	—

Proof-of-concept run, 5,000 steps (Table 5). The two runs differ substantially at step 1 (1.263 vs. 0.723) despite sharing seed 42 and an identical learning-rate schedule—both curves report the same learning rate at every logged step, so warmup does not explain the gap. The cause is the data partition: the distributed run’s first step draws its 512 tokens from the heads of two disjoint 250k-token shards, while the baseline draws from the full 500k-token stream. Different first batches give different initial losses. This means the comparison holds the algorithm and schedule fixed but not the data order, so it is a test of whether gradient averaging converges to the same solution, not a step-for-step equivalence check. Both runs converge to essentially the same final loss (ratio 1.005). The mean loss over all steps is 20% higher for the distributed run because of the early-step gap, not because of slower final convergence.

Full run, 50,000 steps (Table 6). At 50,000 steps, the two-worker SAE matches the single-node baseline to within 0.05% on MSE, 0% on L0 and FVE, and 0.2 percentage points on dead feature rate. The two-worker run is $1.45\times$ slower, but since both run in one process that figure reflects only the sequential stepping of two half-size batches plus the parameter-averaging step—it is not a measurement of communication cost. A real two-node deployment would add link latency and serialisation that this experiment does not model, and could recover some of it by overlapping gradient exchange with the next forward pass. We make no quantitative claim about real distributed overhead.

Figure 1 plots both trajectories. After the first few hundred steps the curves are visually indistinguishable, with loss declining from ~ 0.02 at step 500 to ~ 0.0178 at step 5,000.

Table 6: “Distributed” vs. single-node SAE training at 50,000 steps. **Both runs execute in a single Python process on one machine:** the two workers are stepped sequentially and their parameters averaged, which is algebraically identical to synchronous data-parallel training but exercises none of the inter-process or inter-node machinery. This measures whether the aggregation rule preserves convergence; it does not measure a distributed system. The reported runtime ratio is therefore a lower bound on real distributed overhead.

Metric	Distributed	Single-node	Ratio
Final MSE	0.010613	0.010608	1.0005
Final L0	64.0	64.0	1.0
Final FVE	0.9822	0.9822	1.0
Dead features	1,109 (27.1%)	1,117 (27.3%)	—
Runtime	79.7 min	55.0 min	1.45×
Verdict	PASS: within 0.05% MSE of single-node baseline		

4.3 SAE Quality Metrics

Table 7 reports reconstruction quality for the three 16,384-feature SAEs used in the cross-architecture analysis. Every figure is computed over the full 500,000-token corpus rather than read off the last training batch. This distinction matters more than it might appear: the per-batch **fve** logged during training swings by several points between adjacent steps (Qwen’s last ten logged values span 0.895–0.992), so a final-step reading is a sample from a wide distribution, not a summary of it. Our earlier draft quoted those single-batch values; for Mistral the difference is 3.8 percentage points.

Fraction of Variance Explained (FVE) is the primary metric; raw MSE is activation-scale-dependent and is not comparable across models with different hidden dimensions and activation norms.

Table 7: SAE reconstruction quality over the full 500k-token corpus. All three SAEs are trained to the same specification (50,000 steps, 500k tokens). Dead-feature counts are features that never fire anywhere in the corpus, a stricter and more stable criterion than the rolling 5,000-step **dead_5k** window logged during training. The only genuine held-out measurement comes from a superseded 1,000-step Mistral checkpoint.

Model	Layer	Dict	Steps	MSE	FVE	Dead features	Evaluation
Llama-3.2-3B	14	16,384	50,000	0.00617	0.9895	3,344 (20.4%)	in-sample
Mistral-7B [‡]	16	16,384	50,000	0.00195	0.9659	2,346 (14.3%)	in-sample
Qwen2.5-3B	18	16,384	50,000	0.27532	0.9839	9,398 (57.4%)	in-sample
<i>Superseded 1,000-step Mistral checkpoint, retained for the held-out measurement:</i>							
Mistral-7B ^{‡*}	16	16,384	1,000	0.00415	0.9273	24 (0.1%)	in-sample
Mistral-7B ^{‡*}	16	16,384	1,000	0.00426	0.9252	26 (0.2%)	held-out

[‡] 4-bit quantized model weights. * Superseded checkpoint, trained on the first 50k tokens only; retained because it is the workspace’s only genuine held-out split.




Figure 1: Two-worker (averaged) vs. single-node SAE training loss on Llama-3.2-3B layer-14 activations (dict = 4,096, $k = 64$). Left: full 5,000-step trajectory, log scale. Right: steps 500–5,000 on a linear scale, showing the two runs converging to within 0.5% of each other. The step-1 gap reflects different first batches under the data partition, not a difference in schedule. Data: `data/distributed-sae-training-validation.json`.

None of the converged numbers is held-out. All three SAEs were trained for 205 epochs over their entire 500k-token activation dump. There is no unseen data for any of them: every figure in the top three rows is in-sample, and at that level of repetition an SAE can memorise the corpus rather than learn a dictionary. We report their FVE because it is the best available characterisation, not because it establishes generalisation.

One genuine held-out measurement exists, from a superseded checkpoint. An earlier Mistral run trained on only the first 50,000 tokens, leaving 450,000 unseen; it reconstructs held-out activations at FVE 0.9252 against an in-sample 0.9273, a gap of 0.2 percentage points. Whatever else is wrong with a 1,000-step SAE, it is not overfitting its training corpus. We retain that checkpoint so the measurement stays reproducible, but it comes from a model whose features had barely specialised—a converged SAE has far more capacity to memorise—so it is suggestive rather than probative. Reserving a held-out activation split *before* training is the single cheapest fix available to this work, and none of these runs did it.

Dead features and dictionary collapse. Qwen’s 57.4% dead rate is the dominant quality problem, against 20.4% for Llama and 14.3% for Mistral at the same step count. Our training code implements no auxiliary dead-feature revival loss.

The dynamics behind these numbers are not monotonic, and are worth stating because they are easy to misread from a final checkpoint. All three runs undergo an abrupt collapse shortly after step 2,000, peaking between steps 5,200 and 5,800 at $3.9\times$ to $41\times$ their pre-collapse dead count, followed by a long partial recovery: 57% of collapsed features return in Llama, 56% in Mistral, 16% in Qwen. A count read mid-training therefore overstates permanent loss—an earlier reading of the Llama run at step 13,000 gave 5,278 dead against 3,753 at convergence. Recovery is possible because b_{dec} continues to receive gradient from the features that are firing, so a silent feature’s pre-activation keeps moving even while its

own encoder row is frozen. Adding the auxiliary loss of Gao et al. [5] is a prerequisite for any future run.

Dense features dominate the sparse code. A further pathology cuts across all three SAEs and bears directly on the safety-classifier result in §4.5. With $k = 128$ of 16,384 features active per token, a feature should fire on roughly 0.78% of tokens if activation were spread evenly. Instead, 25 Llama features, 20 Qwen features and 16 Mistral features fire on more than half of all tokens, and the 200 most frequent features absorb 31%, 39% and 29% of total activation mass respectively—1.2% of the dictionary carrying roughly a third of the code. These are dense, low-specificity directions closer to a learned bias than to an interpretable feature. Any procedure that selects features by activation frequency draws heavily from this degenerate head. Training to convergence dilutes the concentration but does not remove it: on the superseded partial checkpoints the top-200 share was 40% (Llama) and 57% (Mistral).

4.4 Cross-Architecture Feature Matching

The chunk-averaged activation matching analysis was run on all three pairwise combinations of the 16,384-feature SAEs (Llama single-node at step 10k, Mistral at step 1k, Qwen at step 50k), using 50,000 evaluation tokens from WikiText-103 divided into 100 chunks of 500 tokens each.

Feature inventory. Not every dictionary slot is a candidate for matching. Table 8 shows how many features in each SAE are usable: alive on the evaluation corpus and with fingerprint support above the 5% floor.

Table 8: Matchable feature inventory over the 500k-token evaluation corpus. “Dead” features never fire; “low support” features fire in fewer than 50 of 1,000 chunks and cannot be matched meaningfully. Only the final column enters the matching analysis.

Model	Dict	Dead on eval	Low support	Matchable
Llama-3.2-3B	16,384	3,344	913	12,127 (74.0%)
Qwen2.5-3B	16,384	9,398	22	6,964 (42.5%)
Mistral-7B	16,384	2,346	5,863	8,175 (49.9%)

All three SAEs are now trained to the same specification—50,000 steps over 500,000 tokens—and between 43% and 74% of each dictionary is usable. The residual attrition has a single cause: none of the runs used an auxiliary dead-feature revival loss, so Qwen ends 57% dead and Mistral carries a large low-support tail. These are imperfect instruments, but they are no longer *undertrained* instruments, which is what previously blocked any conclusion from the matching analysis.

Permutation null (Table 9). The null mean is the number to sit with. Two features drawn from different models with no functional relationship, matched reciprocally over a 1,000-dimensional fingerprint, correlate at 0.59–0.74 on average. The threshold of 0.8 used

Table 9: Chunk-permutation null over 20 replicates, Pearson metric. τ is set to the null 99th percentile. The rightmost column is the number of matches a chunk-shuffled control produces above that threshold, i.e. the expected false-positive count.

Pair	Null mean	Null p99 ($= \tau$)	Null matches/replicate	Expected FP
Llama–Qwen	0.595	0.950	736	7.4
Mistral–Qwen	0.594	0.944	706	7.1
Llama–Mistral	0.744	0.980	872	8.7

in our earlier analysis is at or below that noise floor for *every* pair. The Llama–Mistral pair is the extreme case: its null mean of 0.744 sits within 0.06 of the threshold that analysis used to declare a match, and its calibrated threshold rises to 0.980. Figure 2 shows the relationship directly.




Figure 2: Permutation-null similarity ranges against the two candidate thresholds. For every model pair, the naive fixed threshold $\tau = 0.80$ (dashed) falls *inside* the null range—between the mean and the 99th percentile of similarities produced by chunk-shuffled controls. The calibrated threshold is the null 99th percentile. Match counts at the calibrated threshold are shown at right. Data: `data/sae-analysis/matching-v2/matching-report.json`.

Matches under the corrected protocol (Table 10).

The universality result is negative. Under the corrected protocol, **exactly one feature triple matches across all three models**. Against a null expectation of 0.15 closed triangles per chunk-shuffled replicate, observing one is unremarkable: under a Poisson model with that rate, $\Pr(\geq 1) = 0.14$. Under the cosine metric the null rate is 0.55 and $\Pr(\geq 1) = 0.42$. A single match is what this procedure produces from noise a substantial fraction of the time, and we do not treat it as a finding. The naive protocol’s 3,753 “universal” features do not survive any one of the three corrections, let alone all three.

Pairwise, a residue survives: 15–23 matched pairs per model pair, against 7–9 expected by chance. The enrichment is real but small—1.7–3.2 $\times$, or on the order of 7–16 genuine

Table 10: Cross-architecture feature matches under reciprocal one-to-one matching at the null-calibrated threshold, with the naive protocol’s counts for comparison. Enrichment is observed matches divided by the expected false-positive count from Table 9.

Pair	Naive protocol	Corrected	Expected FP	Enrichment
Llama–Qwen	5,923	16	7.4	2.2×
Mistral–Qwen	8,099	23	7.1	3.2×
Llama–Mistral	12,174	15	8.7	1.7×
All three (universal)	3,753	1	0.15	—

pairs out of 7,000–12,000 matchable features per model, i.e. roughly 0.1%. With counts this low the Poisson uncertainty alone spans a factor of two, and we do not think the pairwise numbers support a quantitative claim about universality rates. What they support is an upper bound: whatever shared representational vocabulary exists between these three models, this method detects it in well under 1% of the dictionary, not the 23% we previously reported.

Venn structure. With one universal feature the three-circle Venn reduces to near-disjoint sets. Under the Pearson metric, 15 Llama features match Qwen only, 14 match Mistral only, one matches both, and 16,354 of 16,384 match neither. We report this rather than a diagram because a Venn figure at these proportions would be illegible. Note that the corrected region counts partition each dictionary exactly—the naive protocol’s did not, summing to 18,258 regions for a 16,384-feature Qwen dictionary, which was itself a signal that the many-to-one matching had broken the set algebra.

What we can and cannot conclude. An earlier version of this analysis ran on partial checkpoints—the Llama SAE at 20% of its step target, the Mistral SAE at under 2% on a ten-times-smaller corpus—and we could not then separate “these models share little” from “these SAEs are too degraded to detect what they share.” That confound is now substantially removed. All three SAEs are trained to the same specification, and the usable fraction of each dictionary rose from 51% to 74% (Llama) and from 33% to 50% (Mistral). The answer did not move: zero universal features became one, against a null that produces at least one 14% of the time.

We still stop short of concluding that these three models share no representational structure, for two reasons that remain. Qwen’s dictionary is 57% dead and Mistral retains a large low-support tail, both consequences of training without an auxiliary revival loss—so a third to a half of each dictionary is still invisible to the analysis. And the fingerprint statistic itself, mean activation over 500-token chunks, discards within-chunk structure: two features responding to the same rare token in different syntactic contexts are indistinguishable from two features responding to a chunk’s overall topic. A negative result from a coarse statistic is weak evidence of absence.

What we can conclude is methodological, and it generalises past this dataset: cross-model feature matching by activation-pattern similarity requires a permutation null. Without

one, the search over a large dictionary manufactures high similarities from nothing, and a threshold chosen by intuition—0.8 is the natural choice, and the one we made—can sit below the noise floor while appearing stringent. For the Llama–Mistral pair the null mean is 0.744, within 0.06 of that threshold. The gap between 3,753 and 1 is almost entirely attributable to that omission, and training the SAEs properly did not close it.

4.5 Safety Classifier Evaluation

We evaluated a feature-based safety classifier built on the Llama-3.2-3B SAE (layer 14, dict_size = 16,384, $k = 128$, trained 10,000 steps on WikiText-103). The classifier labels SAE features by semantic category using regex-based heuristic matching against maximum-activating token contexts, then aggregates “potentially-harmful” feature activations into a scalar harm score.

Table 11: Feature label distribution among the top-200 most frequently activated features on WikiText-103.

Category	Count	Share
factual	173	86.5%
stylistic	24	12.0%
code-related	2	1.0%
potentially-harmful	1	0.5%

Label distribution (top-200 features by activation frequency).

Table 12: Threshold sweep on the safety classifier. Random baseline F1 = 0.500 (balanced dataset).

Threshold	Precision	Recall	F1	FPR	Notes
0.10	0.500	1.000	0.667	1.000	Degenerate: flags everything
0.50	0.476	0.800	0.597	0.880	
0.70	0.488	0.630	0.550	0.660	
0.90	0.580	0.470	0.519	0.340	Best non-degenerate
1.10	0.732	0.300	0.426	0.110	
1.30	1.000	0.180	0.305	0.000	FPR = 0 but recall collapses

Classifier performance on 200-example balanced eval set (Table 12). The best non-degenerate operating point (threshold = 0.90) achieves F1 = 0.519, essentially indistinguishable from the 0.500 random baseline.

Threshold-free discrimination (Table 13). A threshold sweep alone cannot distinguish “the score carries signal but the threshold is wrong” from “the score carries no signal.” Table 13 reports the threshold-free measures, which settle the question.

Table 13: Threshold-free discrimination of the harm score on the 200-example balanced eval set. The AUC confidence interval is a 2,000-replicate bootstrap.

Metric	Value	Chance level
ROC-AUC	0.564	0.500
95% bootstrap CI	[0.484, 0.645]	<i>includes chance</i>
Average precision	0.657	0.500
Cohen’s d	0.238	0.000
Mean harm score, unsafe	0.882	—
Mean harm score, safe	0.802	—

The 95% confidence interval on ROC-AUC includes 0.500. The classifier is therefore not merely weak but *statistically indistinguishable from chance* at $n = 200$: we cannot reject the hypothesis that the harm score carries no information about the label. The separation between class means (0.882 vs. 0.802) corresponds to $d = 0.238$, a small effect by any convention, and the entire distribution of scores spans only 0.17–1.99. No choice of threshold can rescue a score with this little separation, which is the substantive point the sweep in Table 12 obscures.

Root cause analysis. The failure is diagnostic, not accidental:

1. **Corpus mismatch.** WikiText-103 contains almost no explicitly harmful content. The top-200 features by activation frequency represent the model’s vocabulary for encyclopedic prose—factual, stylistic, code. Safety-relevant features are expected to be low-frequency on this corpus.
2. **Frequency bias, and the degenerate head of the distribution.** Selecting features by activation frequency selects for generalist features, not discriminative ones. This is worse than a mild bias: as §4 shows, 25 of the 200 features the labeller inspected fire on more than half of all tokens, and those 200 features carry 31% of the entire sparse code’s activation mass. The labelling pass was therefore applied almost entirely to dense, low-specificity directions—the features least likely to be about anything in particular. Selection should be by *differential* activation on harmful versus safe text, and should additionally exclude features above a density ceiling.
3. **Single-feature collapse.** The one identified harmful feature (feature 15040, a military-history detector firing on 9.9% of WikiText tokens) is too broad (fires on many safe military-history topics) and too narrow (captures weapons in a Wikipedia context, not harmful instructions in a chat context).
4. **Layer placement.** Layer 14 is at the midpoint of the 28-layer stack. Safety-relevant intent representations are more likely to emerge in later layers (20–28) where the residual stream has had time to integrate broader context.

Implications. The result motivates three specific changes: (a) replace the labeling corpus with a safety-targeted corpus such as AdvBench or HarmBench; (b) select features by *differ-*

ential activation frequency (harmful – safe) rather than absolute frequency; (c) target later layers. The harm scoring pipeline (two-pass frequency analysis, context replay, threshold sweep) is reusable—the bottleneck is input data quality, not pipeline logic.

5 Discussion

5.1 The Infrastructure Case

The primary contribution of this work is not any individual experimental result but the demonstration that the full interpretability stack—distributed inference, live activation streaming, online SAE training, cross-architecture feature analysis, and downstream safety evaluation—can run on two consumer nodes connected by a commodity interface. The alternative (cloud GPU clusters, specialized interconnects, large engineering teams) creates a high barrier to entry that concentrates interpretability research at a small number of well-funded organizations. Democratizing the infrastructure is a precondition for democratizing the science.

The Thunderbolt analysis (§4.1) establishes that the link is not a bottleneck under any foreseeable workload. The binding constraints are compute-side: inference speed on Node A and SAE forward-pass throughput on Node B. Future work should focus on these (larger batches, INT8 quantization for the SAE encoder) rather than the interconnect.

5.2 Universality: A Negative Result and What Produced It

Our earlier analysis reported 3,753 features matching across all three model families and read that as support for the Platonic Representation Hypothesis [7]—the conjecture that large networks trained on diverse data converge on a shared statistical model of reality regardless of architecture. Under a corrected protocol the count is zero, and the honest position is that this work provides no evidence bearing on that hypothesis in either direction.

It is worth being precise about why the original number was so large, because the mechanism is general. Cosine similarity between chunk-averaged SAE fingerprints is a comparison of non-negative vectors, which already carries a large positive baseline. The matching procedure then takes, for each feature, the best of $\sim 8,000$ candidates—an extreme-value statistic whose expectation grows with dictionary size and shrinks with fingerprint dimensionality. At the original resolution of 100 chunks, the best-of-8,000 similarity between chunk-shuffled controls had a 99th percentile of 0.98. A threshold of 0.8 looks stringent when read as a correlation and is in fact permissive by a wide margin. Nothing about this is specific to our models or our SAEs; any study matching features across models by activation-pattern similarity over a modest number of context bins will hit it.

The three procedural corrections compound. One-to-one reciprocal matching alone cuts the Llama–Mistral count from 12,174 to 40 at the same threshold, because the original figure was counting the same handful of Mistral features many times over. Requiring the triangle to close removes the chained triples. Calibrating the threshold against the null removes the rest. No single correction accounts for the gap between 3,753 and 0; all three are necessary. What survives is a weak pairwise signal—10–22 matched pairs per model pair against 5–

7 expected by chance—which is consistent with a small shared vocabulary but far too sparse to characterise. Testing the universality hypothesis properly needs SAEs trained to convergence with dead-feature revival on all three models, a held-out evaluation corpus, and a fingerprint statistic with enough resolution that the null sits well below the decision threshold. None of those three conditions holds here.

The Llama–Mistral asymmetry inverts. Under the naive protocol the Llama–Mistral pair had by far the highest count (12,174; nominally 74% coverage), and we had speculated this reflected shared architectural lineage. Under the corrected protocol Llama–Mistral yields the *fewest* matches (15) and the *lowest* enrichment ($1.7\times$) of the three pairs. The reason is visible in the null: the Llama–Mistral noise floor is far higher than the others (mean 0.744 against 0.59), so its calibrated threshold rises to 0.980 and few pairs clear it. The original asymmetry was an artifact of double-counting individual Mistral features, compounded by a pair-specific noise floor that the fixed threshold ignored—not a fact about the models.

5.3 Safety Classifier: An Honest Failure

The safety classifier result is negative, and its honest framing is more valuable than a positive result would have been. A classifier whose ROC-AUC confidence interval spans chance, accompanied by a mechanistic account of why, is a clearer contribution than one reporting $F1 = 0.8$ on a benchmark without revealing its failure modes.

The mechanistic explanation has two parts. The first is a pipeline design error: features were selected by activation frequency on WikiText-103, a criterion orthogonal to the one the task requires (differential activation on harmful versus safe text). The second compounds it: because 1.2% of the dictionary carries 40% of the sparse code’s activation mass, selecting by frequency selects almost entirely from a degenerate head of dense, low-specificity directions. Fixing the corpus without also excluding dense features would likely not be enough.

The threshold sweep, harm-scoring formula, and evaluation harness are reusable. We would not, however, restate the previous draft’s target of $F1 \geq 0.80$ at $FPR \leq 0.10$ as though it were a projection: nothing in these results licenses a quantitative forecast of what a corrected pipeline would achieve.

5.4 Limitations

The two-node system was never built. This is the limitation that most constrains the paper’s framing. Every measurement reported here ran on a single Mac Studio. The protocol validation used two Python subprocesses over a Unix domain socket; the throughput benchmark used TCP loopback with synthetic tensors; the “distributed” SAE training stepped two workers *sequentially in one process* and averaged their parameters. That last case is the weakest: it validates that the aggregation rule preserves convergence, and nothing about inter-process or inter-node behaviour. The Thunderbolt 4 analysis (§4.1) is arithmetic over datasheet bandwidth, not measurement. Latency, jitter, and failure modes of a real inter-machine link remain unmeasured, and the $1.45\times$ overhead we report for two-worker training is a lower bound on what a real deployment would pay.

No held-out evaluation for two of three SAEs. The Llama and Qwen SAEs were trained for 204 epochs over their entire activation dump, so every reconstruction figure for them is in-sample. Only Mistral has genuine held-out data, and only because it was undertrained. This is a data-collection oversight, not a fundamental constraint—the fix is to reserve a token split before training.

Collapsed dictionaries. All three SAEs now train to the same specification, so under-training is no longer a confound—an earlier version of this work ran on partial checkpoints and could not make that claim. What remains is collapse: the trainer implements no auxiliary revival loss, and Qwen ends with 57.4% of its dictionary dead while Mistral retains a large low-support tail. Between 43% and 74% of each dictionary is usable (Table 8), so a third to a half remains invisible to the matching analysis. We can now say the negative universality result is not explained by undertraining; we cannot yet say it is not explained by collapse.

The fingerprint statistic is coarse. Mean activation over 500-token chunks discards within-chunk structure. Two features that respond to the same rare token in different syntactic contexts are indistinguishable from two features that respond to a chunk’s general topic. A token-level or context-level fingerprint would be more discriminative, at higher cost.

Quantization confound. The Mistral-7B activations were extracted from a 4-bit quantized model, due to authentication constraints on the bf16 base weights. Quantization reduces the precision of individual activations and may suppress features that rely on fine-grained magnitude differences. Our earlier draft argued that the high Llama–Mistral match count showed the representational structure survived quantization; since that count was an artifact, no such reassurance is available, and the confound is unresolved.

Safety evaluation set. The 200-example evaluation set was authored and labelled by a single researcher, with no second annotator and therefore no inter-annotator agreement statistic. Its unsafe examples are described as “inspired by” AdvBench, HarmBench, ToxiGen and the Anthropic red-teaming data rather than drawn from them, which means the results are not comparable to any published benchmark number and the set is not reproducible from its sources. At $n = 200$ the bootstrap CI on AUC spans 0.16, wide enough that a moderately effective classifier and a useless one would be hard to distinguish. Future evaluation should draw directly from a public benchmark at substantially larger n .

No held-out data. All three SAEs trained over 205 epochs of their full activation dump, so every converged reconstruction figure is in-sample. The one held-out measurement comes from a superseded 1,000-step checkpoint. Reserving a split before training is the fix.

Corpus. All SAE training and evaluation used WikiText-103, an encyclopedic corpus that is clean and well-studied but unrepresentative of instruction-following or conversational distributions. This is also the direct cause of the safety classifier’s failure: WikiText contains almost no harmful content, so safety-relevant features have little opportunity to appear.

Scale and split point. All results are for models up to 7B parameters; the 34B feasibility claim is arithmetic, not experiment. The system uses a fixed split at one layer, and full-depth capture—streaming every layer, which is what circuit tracing at scale would require—is specified by the protocol but not demonstrated.

5.5 Path to Production

The immediate next step is implementing the system in Swift/MLX rather than Python, using the same wire protocol. Swift provides access to macOS’s native socket APIs, lower-latency memory management, and better MLX integration for zero-copy buffer passing. The protocol spec is wire-format stable (v1.0.0); Swift and Python implementations are directly interoperable.

Beyond that, the natural extension is online SAE training: rather than collecting activations offline and training on stored data, the SAE trainer on Node B processes the activation stream as it arrives, updating weights continuously. This eliminates the two-stage collect→train pipeline and reduces the hyperparameter tuning cycle from hours to minutes.

The universality question is not closed by our negative result—it is untested, because the instruments were inadequate. Settling it requires, in order: SAEs trained to convergence on all three models with an auxiliary dead-feature revival loss; a held-out activation split reserved before training; and a fingerprint statistic with enough resolution that the permutation null sits well below any threshold of interest. Only then does the “universal SAE” experiment become meaningful—training a single dictionary, initialised from matched features, jointly on all three models’ activations, and asking whether it achieves comparable FVE on each.

6 Conclusion

We have described MLXMESH, a system for distributed mechanistic interpretability on Apple Silicon hardware. The system streams intermediate layer activations between nodes over a purpose-built binary wire protocol, enabling SAE training and analysis without staging activations to disk. We validated the pieces we could validate on one machine: a bit-exact split boundary on Llama-3.2-3B, 1.25 GB/s (10.0 Gbps) protocol throughput over loopback, and two-worker gradient averaging that reaches within 0.05% of the single-node baseline’s MSE. We showed by direct calculation that a Thunderbolt 4 link would carry full-depth activation capture for every model size we consider at under 1.4% utilisation. The two-node deployment itself remains unbuilt, and we have marked every projected figure as projected.

The scientific results are two negatives, reported as such. Cross-architecture feature matching, done properly, finds essentially *no* features shared across Llama-3.2-3B, Mistral-7B and Qwen2.5-3B: one closed triangle, against a null that produces at least one 14% of the time. This corrects our own earlier claim of 3,753, which was an artifact of many-to-one matching, unclosed triangles, and a similarity threshold below the permutation-null mean. The corrected measurement uses SAEs trained to full specification on all three models, so it is not explained by undertraining. A feature-based safety classifier achieves ROC-AUC 0.564 with a 95% confidence interval spanning chance, and we trace the failure to feature selection by

raw activation frequency over a corpus containing almost no harmful content, compounded by a dictionary in which 1.2% of features carry roughly a third of the activation mass.

We think the negatives are the more useful contribution. The infrastructure claim—that two consumer nodes and a commodity interface suffice for interpretability research at this scale—is one we still believe and have partially demonstrated. But the universality result is a cautionary case with a clean moral: when a matching procedure searches thousands of candidates, the similarity threshold must be calibrated against a null, or the procedure will manufacture structure from noise at any scale you run it. We reported 23% universality; the figure detectable by this method, with all three SAEs trained to specification, is indistinguishable from zero.

Acknowledgments

This work was conducted independently using consumer hardware. We thank the developers of MLX, the Hugging Face community for open model weights, and the mechanistic interpretability community whose open publication of methods and findings made this work possible.

Data Availability

All experimental data, training logs, benchmark results, and analysis scripts produced in this work are available at <https://github.com/jmelton-research/mlxmesh>. This includes: activation collection scripts and metadata; SAE training checkpoints and training logs for all four runs; the cross-architecture matching results, including both the superseded v1 outputs and the corrected v2 outputs with permutation nulls (`data/sae-analysis/matching-v2/`); bandwidth benchmark results (clearly labeled as using synthetic activations); the distributed-capture validation JSON; and the safety classifier evaluation dataset and results. The mlxMesh wire protocol specification is documented in `docs/activation-streaming-protocol.md`.

References

- [1] Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Zac Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E. Burke, Tristan Hume, Shan Carter, Tom Henighan, and Christopher Olah. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023. URL <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- [2] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. *arXiv*, 2022. URL <https://arxiv.org/abs/2212.03827>.

- [3] Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. *arXiv*, 2023. URL <https://arxiv.org/abs/2309.08600>.
- [4] Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse, Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher Olah. Toy models of superposition. *Transformer Circuits Thread*, 2022. URL https://transformer-circuits.pub/2022/toy_model/index.html.
- [5] Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders. *arXiv*, 2024. URL <https://arxiv.org/abs/2406.04093>.
- [6] Yanping Huang, Youlong Cheng, Ankur Bapna, Orhan Firat, Mia Xu Chen, Dehao Chen, Hyoungho Lee, Jiquan Ngiam, Quoc V. Le, Yonghui Wu, and Zhifeng Chen. GPipe: Efficient training of giant neural networks using pipeline parallelism. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [7] Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. The platonic representation hypothesis. *arXiv*, 2024. URL <https://arxiv.org/abs/2405.07987>.
- [8] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2019.
- [9] Nikolaus Kriegeskorte, Marieke Mur, and Peter Bandettini. Representational similarity analysis — connecting the branches of systems neuroscience. *Frontiers in Systems Neuroscience*, 2:4, 2008. doi: 10.3389/neuro.06.004.2008.
- [10] Tom Lieberum, Senthooan Rajamanoharan, Arthur Conmy, Lewis Smith, Nicolas Sonnerat, Vikrant Kramár, Neel Nanda, Caglar Sezener, Borja Balle, and Rohin Shah. Gemma Scope: Open sparse autoencoders everywhere all at once on Gemma 2. *arXiv*, 2024. URL <https://arxiv.org/abs/2408.05147>.
- [11] Jack Lindsey, Amanda Gould, Adly Templeton, Callum McDougall, Joshua Batson, Adam Jermyn, Neil Thomas, Carson Denison, Amanda Askill, Anna Durbin, Trenton Bricken, Jason Marcus, Robert Lasenby, Tom Conerly, Tom Henighan, Nick Turner, Brian Chen, David Pearce, Joshua McDonald, Alex Tamkin, Aryan Guha, Aaron Jones, and Christopher Olah. On the biology of a large language model. *Transformer Circuits Thread*, 2025. URL <https://transformer-circuits.pub/2025/attribution-graphs/biology.html>.
- [12] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [13] Deepak Narayanan, Aaron Harlap, Amar Phanishayee, Vivek Seshadri, Nikhil R. Devanur, Gregory R. Ganger, Phillip B. Gibbons, and Matei Zaharia. PipeDream: Generalized pipeline parallelism for DNN training. In *Proceedings of the ACM Symposium on Operating Systems Principles (SOSP)*, 2019.

- [14] Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova Das-Sarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Scott Johnston, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Christopher Olah. In-context learning and induction heads. *Transformer Circuits Thread*, 2022. URL <https://transformer-circuits.pub/2022/in-context-learning-and-induction-heads/index.html>.
- [15] Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, Brian Chen, Adam Pearce, Craig Citro, Emmanuel Scheirer, Andy Jones, Andy Hyun, Alex Tamkin, David Bau, Jack Chen, and Christopher Olah. Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet. *Transformer Circuits Thread*, 2024. URL <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html>.
- [16] Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: A circuit for indirect object identification in GPT-2 small. *arXiv*, 2022. URL <https://arxiv.org/abs/2211.00593>.
- [17] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv*, 2023. URL <https://arxiv.org/abs/2310.01405>.